

# Exam TI / GPU

⌚ Cours

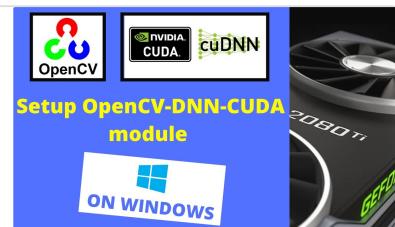
GPU + Traitement d'images

## Installer la version d'OpenCV compilé avec le support de CUDA

Setup OpenCV-DNN module with CUDA backend support on Windows

This video shows step by step tutorial on how to set up the OpenCV-DNN module with CUDA backend support on Windows.

▶ [https://www.youtube.com/watch?v=VAoNMICm\\_Cg](https://www.youtube.com/watch?v=VAoNMICm_Cg)



Setup OpenCV-DNN module with CUDA backend support (For Windows)

In this tutorial, we will be building OpenCV from source with CUDA backend support (OpenCV-DNN-CUDA module).

🎥 <https://medium.com/geekculture/setup-opencv-dnn-module-with-cuda-backend-support-for-windows-7f1856691da3>



Install CUDA and CUDNN on Windows & Linux

CONTENTS

🎥 <https://medium.com/geekculture/install-cuda-and-cudnn-on-windows-linux-52d1501a8805>



### Vérification si tout est bien installé

```
import numpy as np
import cv2 as cv
import time
```

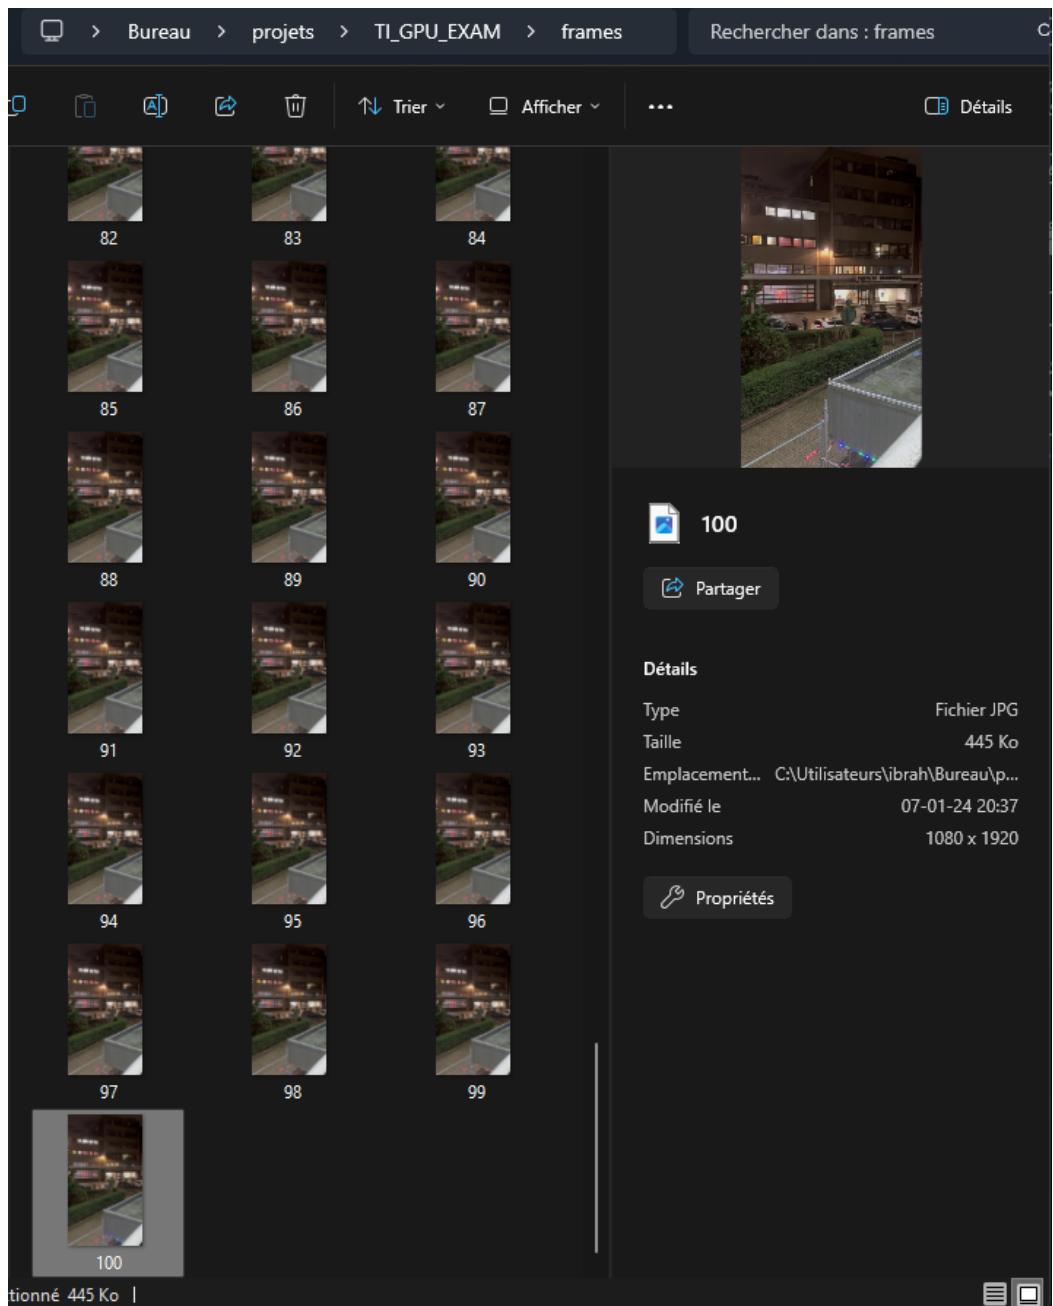
```
npTmp = np.random.random((1024, 1024)).astype(np.float32)
npMat1 = np.stack([npTmp, npTmp], axis=2)
npMat2 = npMat1
cuMat1 = cv.cuda_GpuMat()
cuMat2 = cv.cuda_GpuMat()
cuMat1.upload(npMat1)
cuMat2.upload(npMat2)
start_time = time.time()
cv.cuda.gemm(cuMat1, cuMat2, 1, None, 0, None, 1)

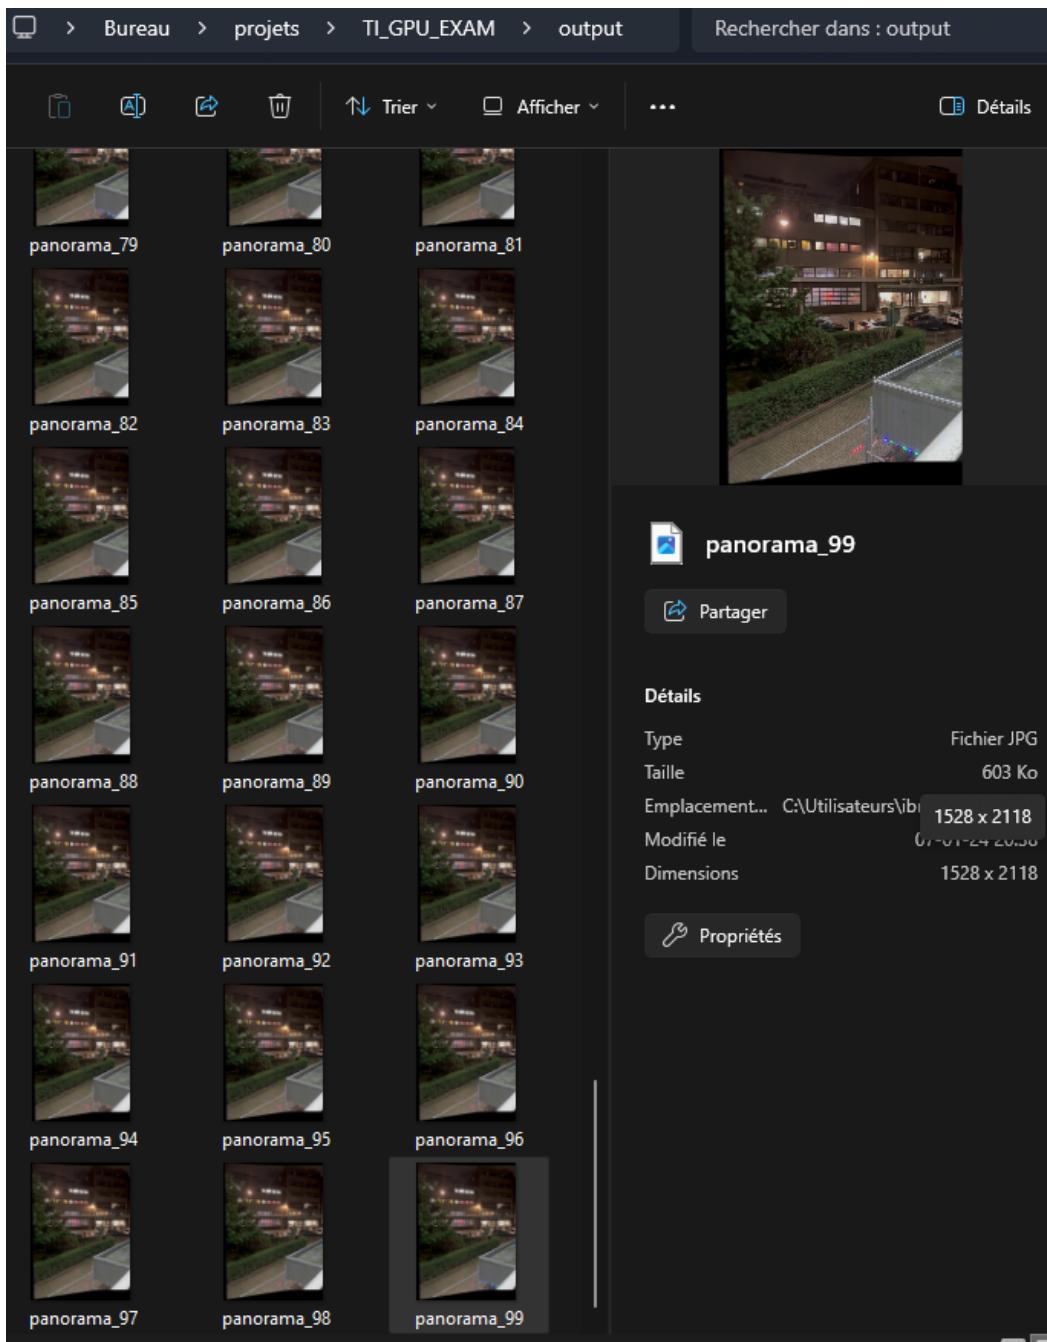
print("CUDA using GPU --- %s seconds ---" % (time.time() - start_time))
start_time = time.time()
cv.gemm(npMat1, npMat2, 1, None, 0, None, 1)
print("CPU --- %s seconds ---" % (time.time() - start_time))
```

output :

```
(env_gpu) PS C:\Users\ibrah\Desktop\projets\Traitement images> python
CUDA using GPU --- 0.11492133140563965 seconds ---
CPU --- 1.7428045272827148 seconds ---
```

**===== ==Test pour 100 frames CPU ======**





## Profilage de performance sans CUDA

### The Python Profilers

Source code: Lib/profile.py and Lib/pstats.py  
Introduction to the profilers:  
cProfile and profile provide deterministic profiling of Python programs. A profile is a set of statistics that describes...

 <https://docs.python.org/3/library/profile.html>



Pour commencer je fais un profilage de performance pour comprendre où le programme passe plus de temps lors de son exécution, ce qui me permet d'identifier les parties de mon code pouvant être optimisé.

```
(env_gpu) PS C:\Users\ibrah\Desktop\projets\exam_Ti_gpu> python .\app.py
103801 function calls in 13.633 seconds

Ordered by: internal time

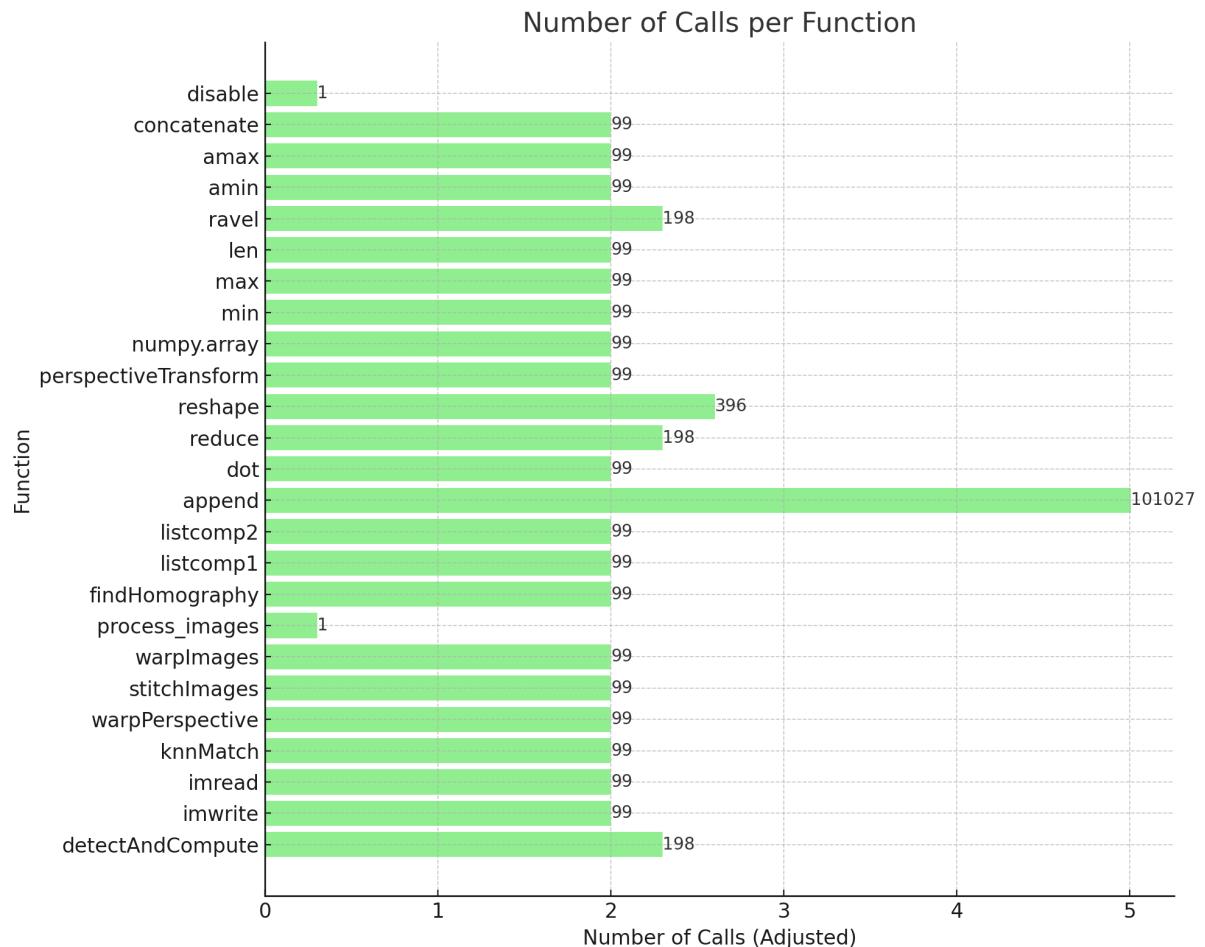
      ncalls  tottime  percall  cumtime  percall   filename:lineno(function)
          198    5.421    0.027    5.421    0.027 {method 'detectAndCompute' of 'cv2.Feature2D' objects}
           99    3.938    0.040    3.938    0.040 {imwrite}
           99    2.115    0.021    2.115    0.021 {imread}
           99    1.106    0.011    1.106    0.011 {method 'knnMatch' of 'cv2.DescriptorMatcher' objects}
           99    0.605    0.006    0.605    0.006 {warpPerspective}
           99    0.130    0.001    7.498    0.076 app.py:71(stitchImages)
           99    0.084    0.001    0.699    0.007 app.py:44(warpImages)
            1    0.082    0.082  13.633   13.633 app.py:96(process_images)
           99    0.071    0.001    0.071    0.001 {findHomography}
           99    0.042    0.000    0.042    0.000 app.py:84(<listcomp>)
           99    0.019    0.000    0.019    0.000 app.py:85(<listcomp>)
          101027    0.010    0.000    0.010    0.000 {method 'append' of 'list' objects}
           99    0.006    0.000    0.006    0.000 {method 'dot' of 'numpy.ndarray' objects}
          198    0.002    0.000    0.002    0.000 {method 'reduce' of 'numpy.ufunc' objects}
          396    0.000    0.000    0.000    0.000 {method 'reshape' of 'numpy.ndarray' objects}
           99    0.000    0.000    0.000    0.000 {perspectiveTransform}
           99    0.000    0.000    0.000    0.000 {built-in method numpy.array}
           99    0.000    0.000    0.002    0.000 {method 'min' of 'numpy.ndarray' objects}
           99    0.000    0.000    0.001    0.000 {method 'max' of 'numpy.ndarray' objects}
           99    0.000    0.000    0.000    0.000 {built-in method builtins.len}
          198    0.000    0.000    0.000    0.000 {method 'ravel' of 'numpy.ndarray' objects}
           99    0.000    0.000    0.001    0.000 _methods.py:43(_amin)
           99    0.000    0.000    0.001    0.000 _methods.py:39(_amax)
           99    0.000    0.000    0.000    0.000 multiarray.py:153(concatenate)
            1    0.000    0.000    0.000    0.000 {method 'disable' of '_lsprof.Profiler' objects}

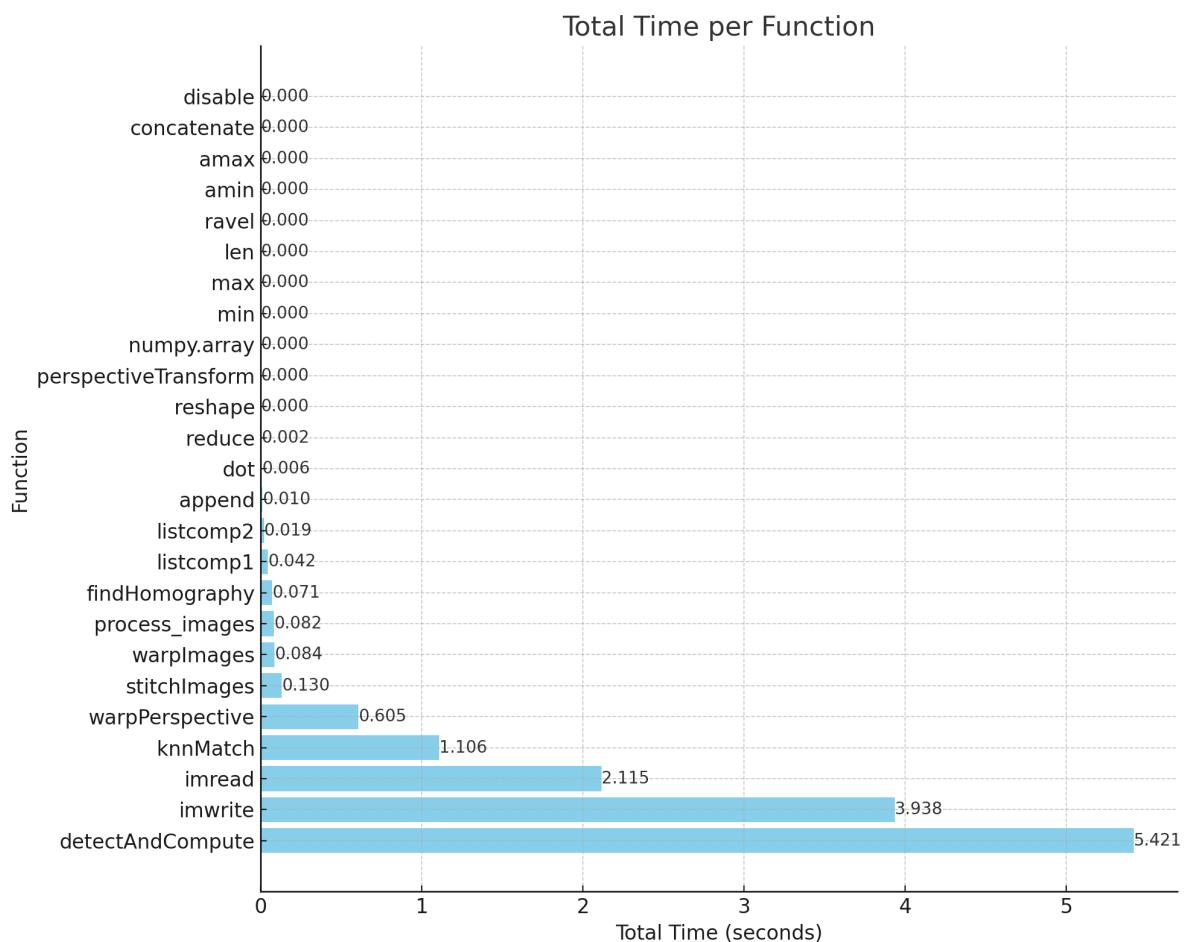
Temps d'exécution sans CUDA: 27.33 secondes
```

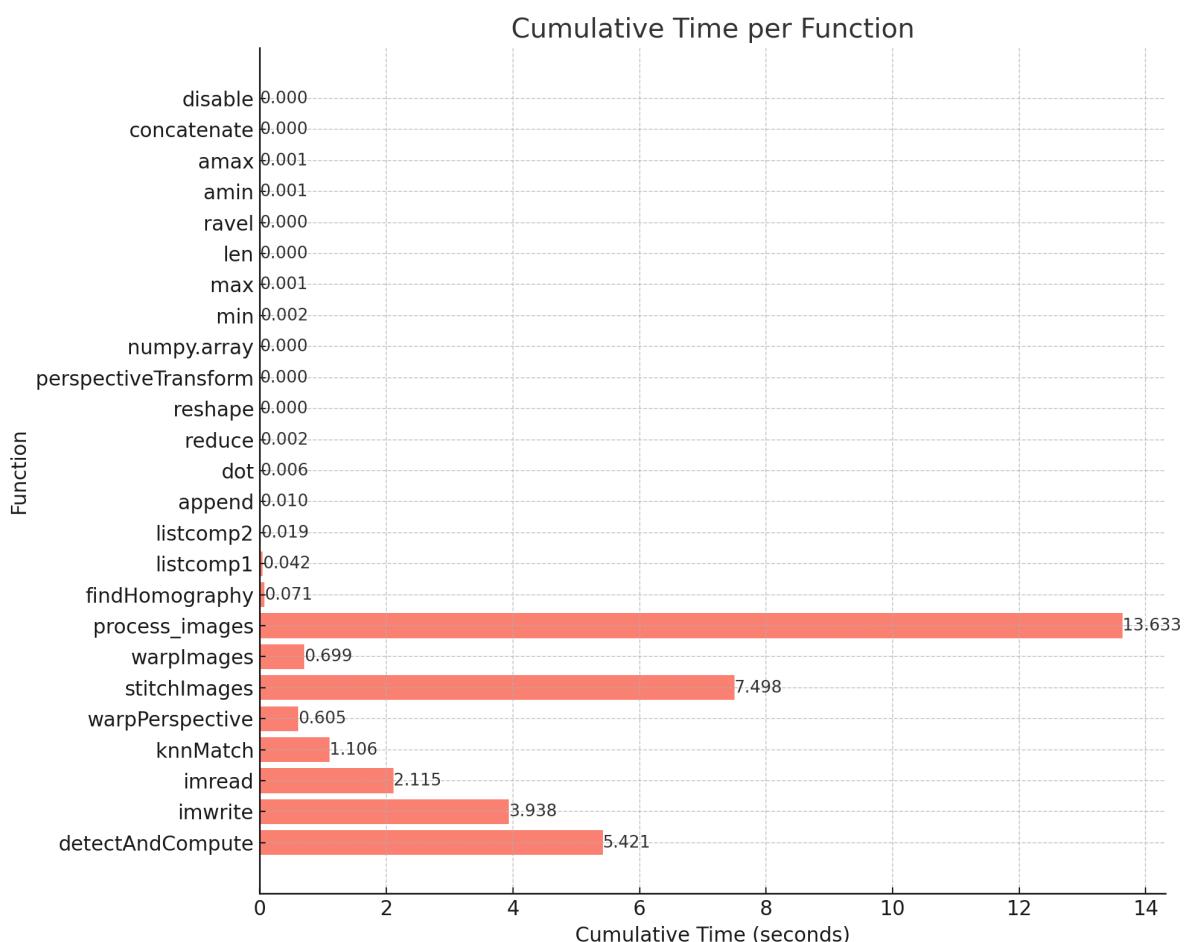
Le résultat est le profil de performance du script Python. Il donne des informations détaillées sur le temps d'exécution de chaque fonction dans le script.

- `ncalls` : le nombre d'appels à la fonction.
- `tottime` : le temps total passé dans la fonction donnée (sans compter le temps passé dans les appels à d'autres fonctions).
- `percall` : le temps moyen passé dans la fonction par appel (tottime/ncalls).
- `cumtime` : le temps cumulé passé dans la fonction donnée, y compris le temps passé dans tous les appels de fonctions à partir de cette fonction.
- `percall` : le temps moyen passé dans la fonction par appel (cumtime/ncalls).
- `filename:lineno(function)` : le nom du fichier, le numéro de ligne et le nom de la fonction.

Dans ce cas, la fonction `detectAndCompute` de l'objet `cv2.Feature2D` est celle qui prend le plus de temps (`tottime`), suivie par `imwrite`, `imread`, `knnMatch` de `cv2.DescriptorMatcher` et `warpPerspective`.








---

## CPU optimisé imread et imwrite

---

OUTPUT

```
(env_gpu) PS C:\Users\ibrah\Desktop\projets\exam_Ti_gpu> python .\ap
103793 function calls in 7.505 seconds
```

Ordered by: internal time

ncalls	tottime	percall	cumtime	percall	filename:lineno(function)
198	5.379	0.027	5.379	0.027	{method 'detectAndComp
99	1.058	0.011	1.058	0.011	{method 'knnMatch' of
99	0.588	0.006	0.588	0.006	{warpPerspective}
99	0.127	0.001	7.388	0.075	app.py:69(stitchImages)
99	0.083	0.001	0.681	0.007	app.py:42(warpImages)
99	0.072	0.001	0.072	0.001	{findHomography}
1	0.062	0.062	7.505	7.505	app.py:95(process_imag
1	0.055	0.055	0.055	0.055	{imwrite}
99	0.041	0.000	0.041	0.000	app.py:82(<listcomp>)
99	0.019	0.000	0.019	0.000	app.py:83(<listcomp>)
101215	0.010	0.000	0.010	0.000	{method 'append' of 'l
99	0.006	0.000	0.006	0.000	{method 'dot' of 'nump
198	0.002	0.000	0.002	0.000	{method 'reduce' of 'n
396	0.000	0.000	0.000	0.000	{method 'reshape' of '
99	0.000	0.000	0.000	0.000	{perspectiveTransform}
99	0.000	0.000	0.000	0.000	{built-in method numpy
99	0.000	0.000	0.002	0.000	{method 'min' of 'nump
99	0.000	0.000	0.001	0.000	{method 'max' of 'nump
100	0.000	0.000	0.000	0.000	{built-in method built
198	0.000	0.000	0.000	0.000	{method 'ravel' of 'nu
99	0.000	0.000	0.002	0.000	_methods.py:43(_amin)
99	0.000	0.000	0.001	0.000	_methods.py:39(_amax)
99	0.000	0.000	0.000	0.000	multiarray.py:153(concat
1	0.000	0.000	0.000	0.000	{method 'disable' of '

Temps d'exécution sans CUDA: 18.13 secondes

Sun Jan 7 04:04:59 2024

```
+-----+-----+-----+
| NVIDIA-SMI 546.01          Driver Version: 546.01      CUD
+-----+-----+-----+
| GPU  Name                  TCC/WDDM    | Bus-Id      Disp.A | V
| Fan  Temp     Perf        Pwr:Usage/Cap |             Memory-Usage | G
|          |          |          |          |          |          |
|-----+-----+-----+-----+-----+-----+-----+-----+
|  0  NVIDIA GeForce RTX 2070 ... WDDM    | 00000000:01:00.0  On | | | | |
| N/A   75C     P8           17W /  83W |      646MiB /  8192MiB |
|          |          |          |          |          |          |
+-----+-----+-----+-----+-----+-----+-----+
```

Processes:						
GPU	GI	CI	PID	Type	Process name	
	ID	ID				
<hr/>						
0	N/A	N/A	1280	C+G	...t.LockApp_cw5n1h2txyewy\Lock	
0	N/A	N/A	5064	C+G	... Access Service\PowerButton	
0	N/A	N/A	10068	C+G	C:\Windows\explorer.exe	
0	N/A	N/A	10540	C+G	...CBS_cw5n1h2txyewy\TextInputH	
0	N/A	N/A	13540	C+G	...nt.CBS_cw5n1h2txyewy\SearchH	
0	N/A	N/A	13648	C+G	...2txyewy\StartMenuExperienceH	
0	N/A	N/A	18732	C+G	...B\system_tray\lghub_system_t	
0	N/A	N/A	19076	C+G	C:\Program Files\NordVPN\NordVP	
0	N/A	N/A	19440	C+G	...inaries\Win64\EpicGamesLaunc	
0	N/A	N/A	23812	C+G	...5n1h2txyewy\ShellExperienceH	
0	N/A	N/A	41752	C+G	...crosoft\Edge\Application\mse	
0	N/A	N/A	60704	C+G	...ekyb3d8bbwe\PhoneExperienceH	
0	N/A	N/A	78956	C+G	...GeForce Experience\NVIDIA Sh	
0	N/A	N/A	80756	C+G	...wekyb3d8bbwe\XboxGameBarWidge	
0	N/A	N/A	85456	C	...h\anaconda3\envs\env_gpu\pyt	

---

(env\_gpu) PS C:\Users\ibrah\Desktop\projets\exam\_Ti\_gpu> python .\ap  
103793 function calls in 8.069 seconds

Ordered by: internal time

ncalls	tottime	percall	cumtime	percall	filename:lineno(function)
198	5.865	0.030	5.865	0.030	{method 'detectAndComp
99	1.103	0.011	1.103	0.011	{method 'knnMatch' of
99	0.611	0.006	0.611	0.006	warpPerspective}
99	0.131	0.001	7.954	0.080	app.py:69(stitchImages)
99	0.090	0.001	0.710	0.007	app.py:42(warpImages)
99	0.072	0.001	0.072	0.001	{findHomography}
1	0.064	0.064	8.069	8.069	app.py:95(process_imag
1	0.050	0.050	0.050	0.050	{imwrite}
99	0.042	0.000	0.042	0.000	app.py:82(<listcomp>)
99	0.020	0.000	0.020	0.000	app.py:83(<listcomp>)
101215	0.010	0.000	0.010	0.000	{method 'append' of 'l
99	0.006	0.000	0.006	0.000	{method 'dot' of 'nump

```

198    0.002    0.000    0.002    0.000 {method 'reduce' of 'n
396    0.000    0.000    0.000    0.000 {method 'reshape' of 'n
99     0.000    0.000    0.000    0.000 {perspectiveTransform}
99     0.000    0.000    0.000    0.000 {built-in method numpy
99     0.000    0.000    0.002    0.000 {method 'min' of 'nump
99     0.000    0.000    0.001    0.000 {method 'max' of 'nump
100    0.000    0.000    0.000    0.000 {built-in method built
198    0.000    0.000    0.000    0.000 {method 'ravel' of 'nu
99     0.000    0.000    0.002    0.000 _methods.py:43(_amin)
99     0.000    0.000    0.001    0.000 _methods.py:39(_amax)
99     0.000    0.000    0.000    0.000 multiarray.py:153(concat
1      0.000    0.000    0.000    0.000 {method 'disable' of '
```

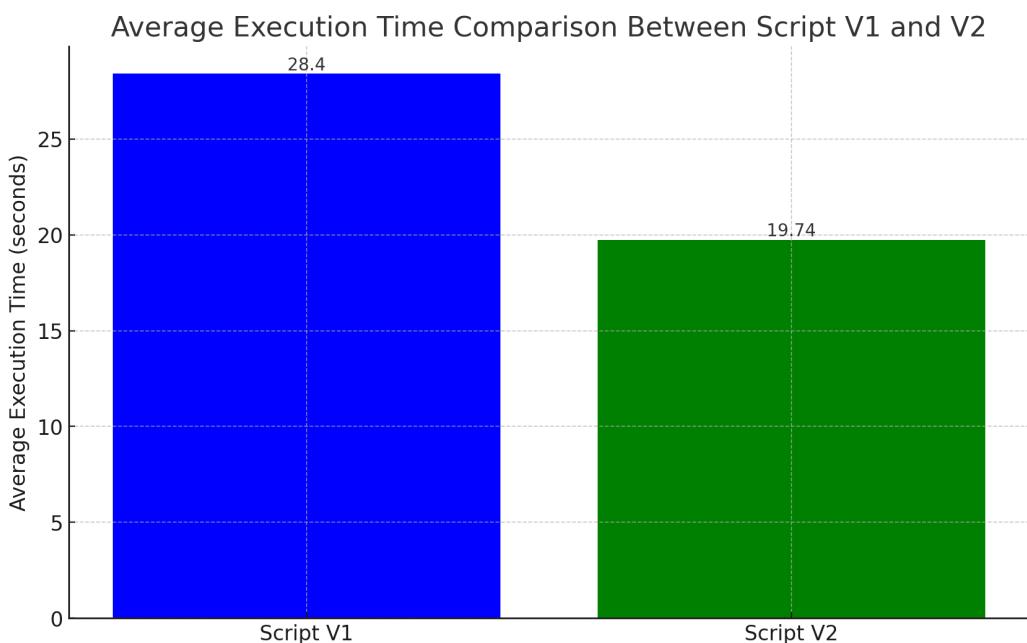
Temps d'exécution sans CUDA: 18.68 secondes

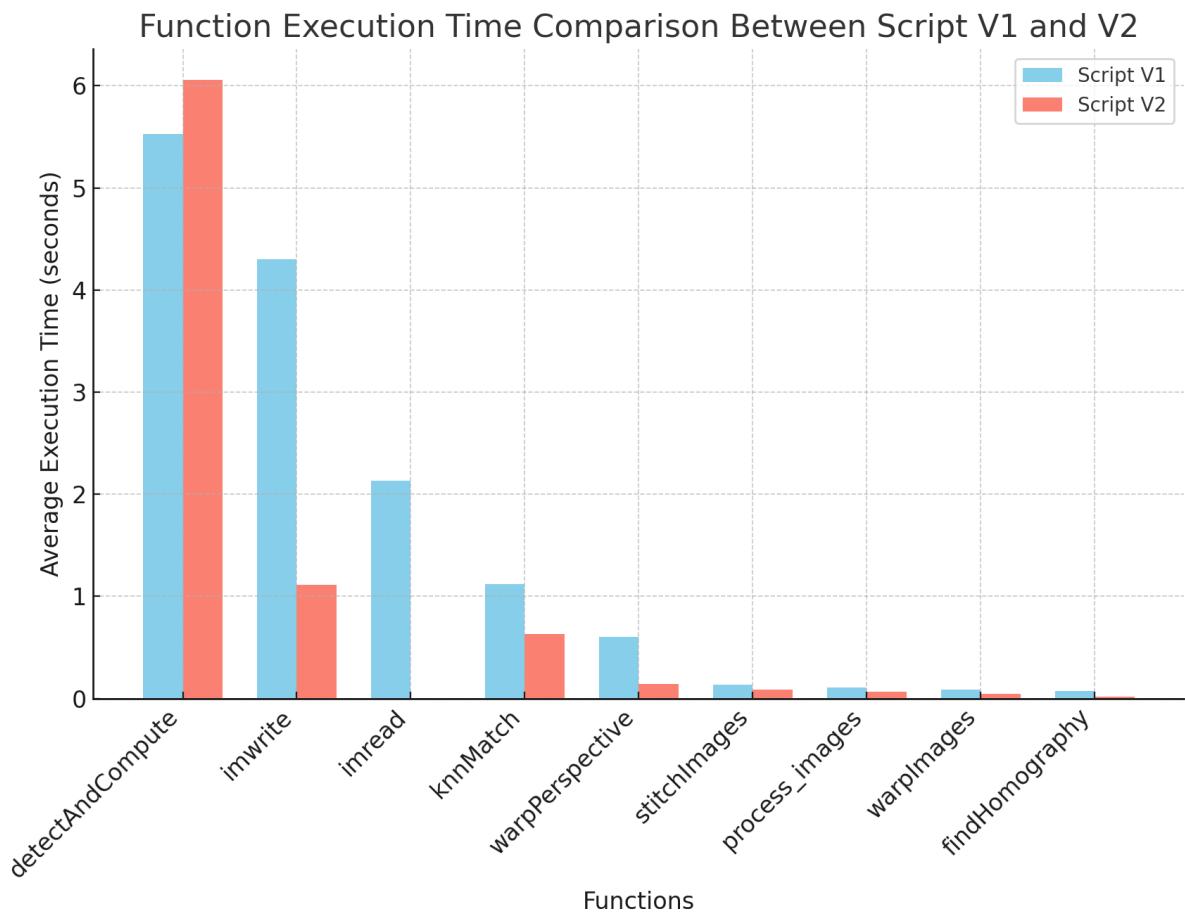
Sun Jan 7 04:05:25 2024

```
+
| NVIDIA-SMI 546.01                 Driver Version: 546.01       CUD
+-----+-----+-----+
| GPU  Name                   TCC/WDDM | Bus-Id      Disp.A | V
| Fan  Temp     Perf          Pwr:Usage/Cap |           Memory-Usage | G
|          |          |          |          |          |          |
|-----+-----+-----+-----+-----+-----+
|   0  NVIDIA GeForce RTX 2070 ... WDDM  | 00000000:01:00.0  On | | | | |
| N/A   76C     P8             9W /  81W |      646MiB /  8192MiB |
|          |          |          |          |          |          |
+-----+-----+-----+-----+-----+-----+
+
| Processes:
| GPU  GI  CI          PID  Type  Process name
| ID   ID
|-----+-----+-----+-----+-----+-----+
|   0  N/A  N/A        1280  C+G  ...t.LockApp_cw5n1h2txyewy\Lock
|   0  N/A  N/A        5064  C+G  ... Access Service\PowerButton
|   0  N/A  N/A       10068  C+G  C:\Windows\explorer.exe
|   0  N/A  N/A       10540  C+G  ...CBS_cw5n1h2txyewy\TextInputH
|   0  N/A  N/A       13540  C+G  ...nt.CBS_cw5n1h2txyewy\SearchH
|   0  N/A  N/A       13648  C+G  ...2txyewy\StartMenuExperienceH
|   0  N/A  N/A       18732  C+G  ...B\system_tray\lghub_system_t
|   0  N/A  N/A       19076  C+G  C:\Program Files\NordVPN\NordVP
|   0  N/A  N/A       19440  C+G  ...inaries\Win64\EpicGamesLaunc
```

	0	N/A	N/A	23812	C+G	...5n1h2txyewy\ShellExperienceH
	0	N/A	N/A	41752	C+G	...crosoft\Edge\Application\mse
	0	N/A	N/A	60704	C+G	...ekyb3d8bbwe\PhoneExperienceH
	0	N/A	N/A	78956	C+G	...GeForce Experience\NVIDIA Sh
	0	N/A	N/A	80756	C+G	...wekyb3d8bbwe\XboxGameBarWidg
	0	N/A	N/A	82764	C	...h\anaconda3\envs\env_gpu\pyt
+						

(env\_gpu) PS C:\Users\ibrah\Desktop\projets\exam\_Ti\_gpu>





## Algorithme de detection ORB

**detectAndCompute** → Keypoint et Descripteurs

Pour detecter les kp et Dscp

Kp = mini region depixel unique / intéressant —→ en `print(keypoints[0])` on a ca : `<KeyPoint 0x7f8e6c8>`

- **kp:** Ce sont des points spécifiques dans une image qui sont uniques et peuvent être utilisés pour comparer différentes images. Par exemple, les coins, les bords et les taches sont souvent utilisés comme points d'intérêt car ils sont uniques et peuvent être facilement identifiés dans différentes images.

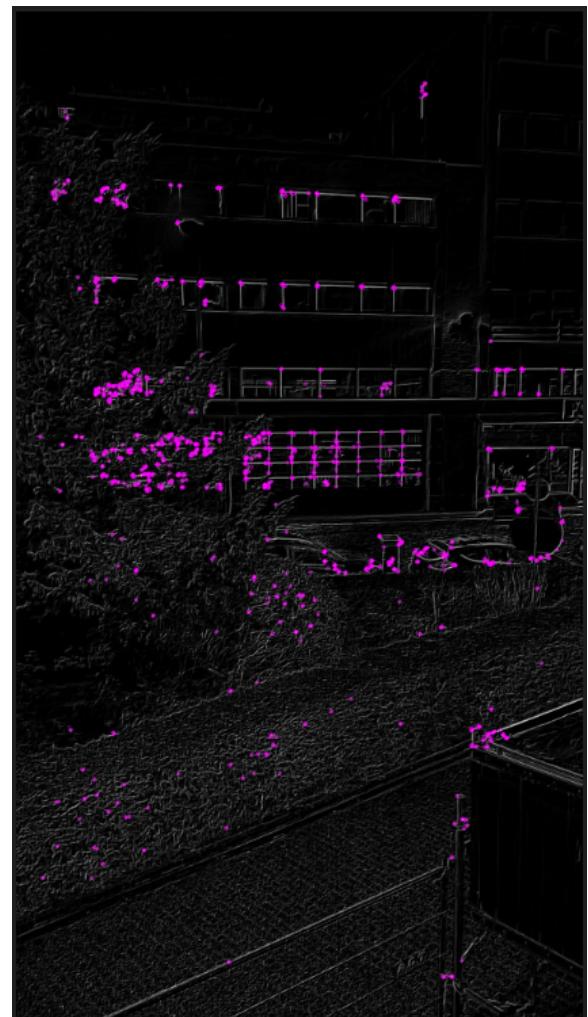
il contiennent les info comme la position du point d'intérêt (x, y), la taille de la région d'intérêt

- **Dscp** : est représentation numérique d'un kp et donc sert à comparer les points entre 2 img

```
print(descriptors[0]) [ 49 241 236 253 255 159 254 223 63 254 255 254 255 127 191  
251 255 255 255 255 255 255 255 255 255 255 255 255 255 255 255 255 255]
```

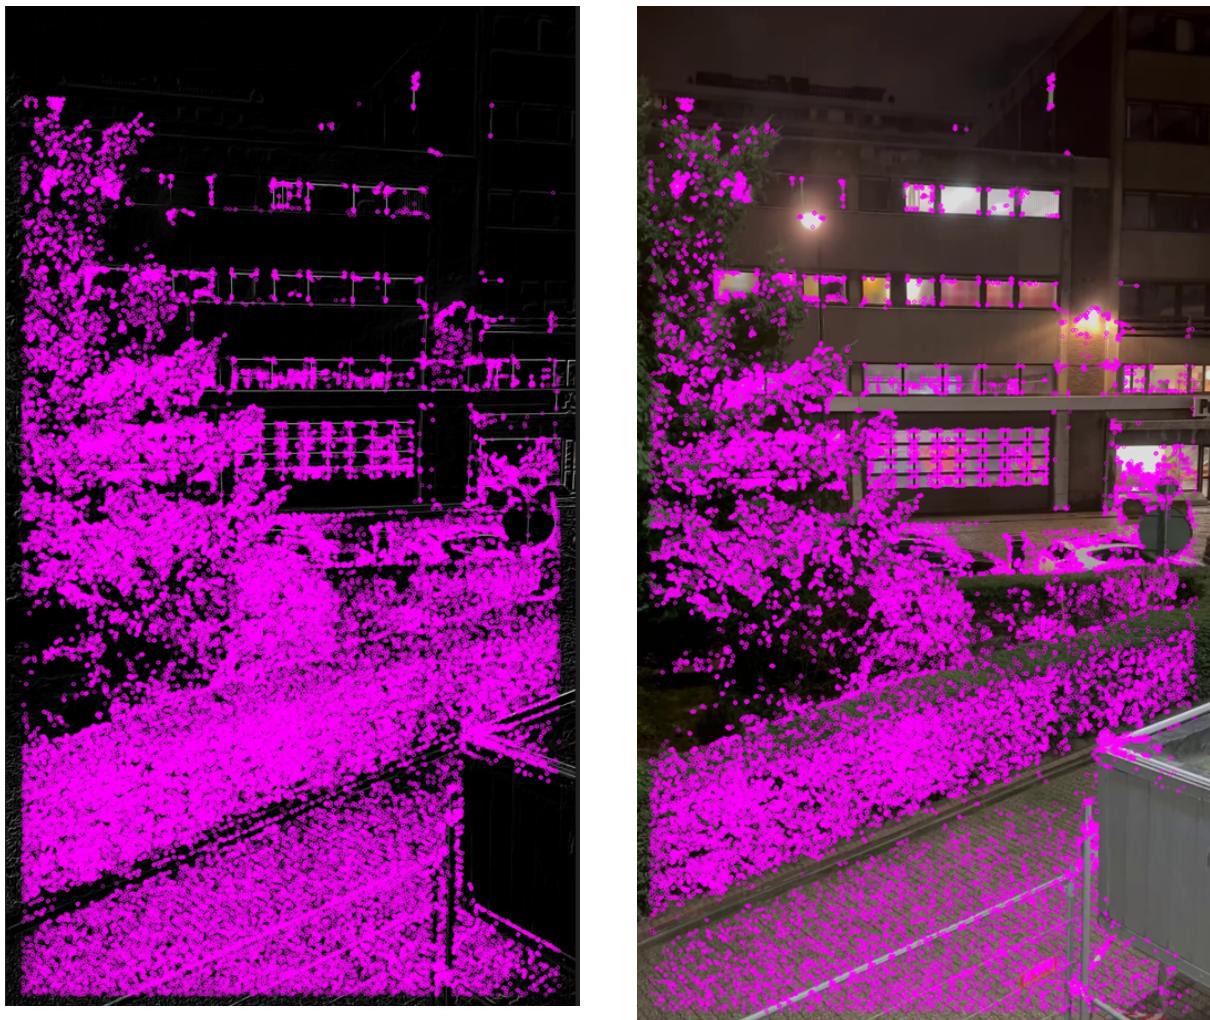
Pour `orb = cv2.ORB_create(nfeatures=2000)`

Temps d'exécution : 0.35 secondes en moyenne



A 200000

Temps d'exécution : 0.56 secondes



## **UTILISATION DE CUDA POUR ENCORE OPTIMISER LE SCRIPT GRACE A LA PUISANCE DU GPU**

**Filtre de Sobel pour aider les algo de detection  
ici ORB**

Il aide à mieux détecter les bordures d'images ce qui peut rendre l'alignement des images pour la construction du panorama plus précis

## Avec filtre sobel



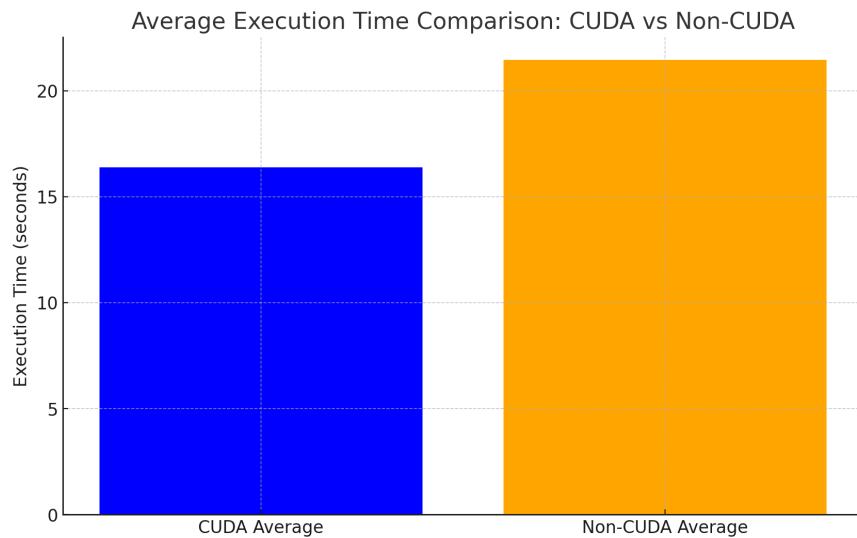
## Sans filtre sobel

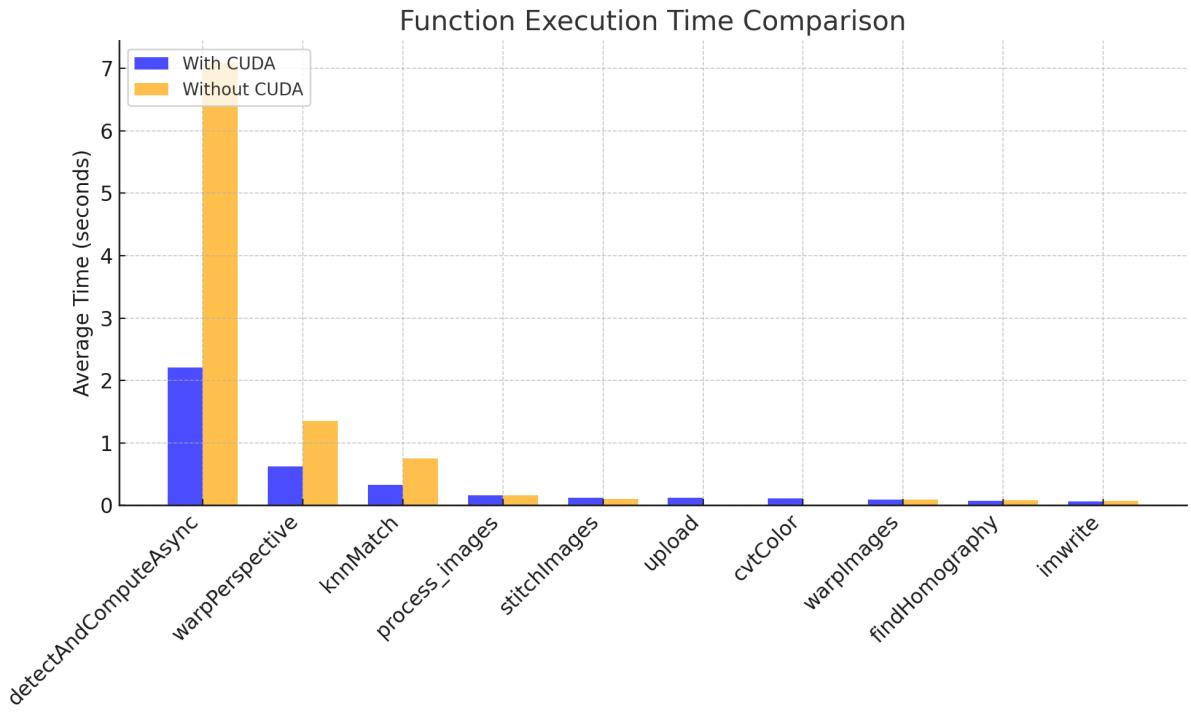


**detectAndCompute** faire tout cet ensemble en gpu et renvoyer les resultat au cpu

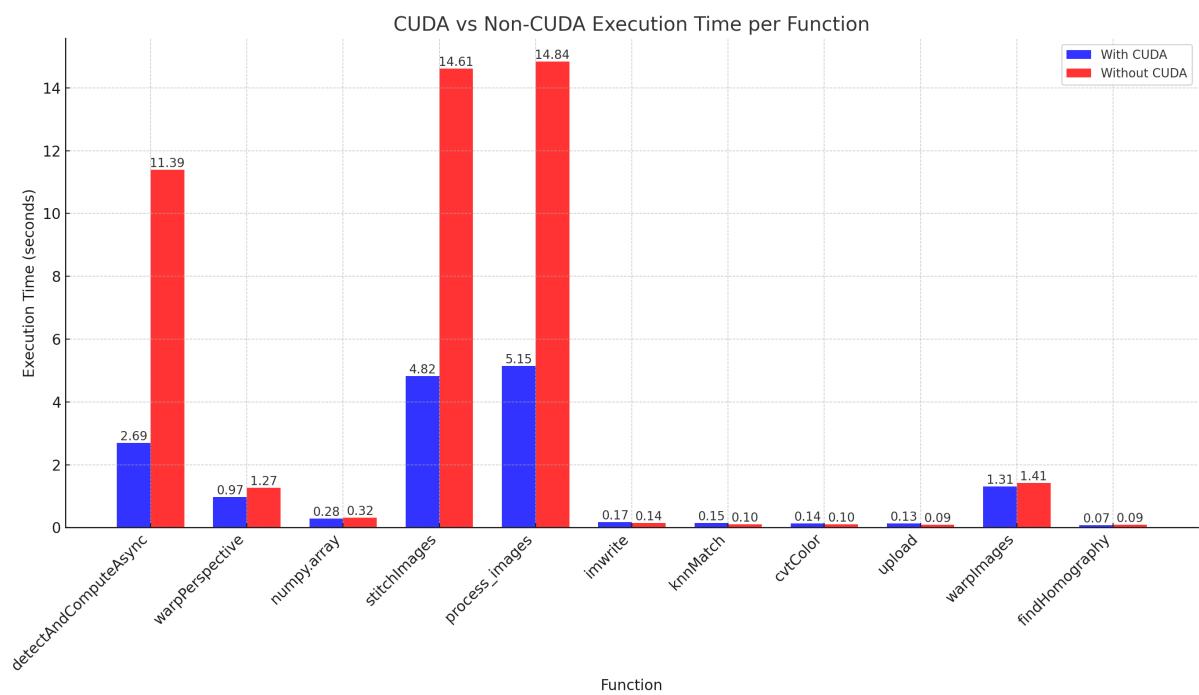
`matches = bf.knnMatch(descriptors1, descriptors2, k=2)` envoie directement le resultat au cpu

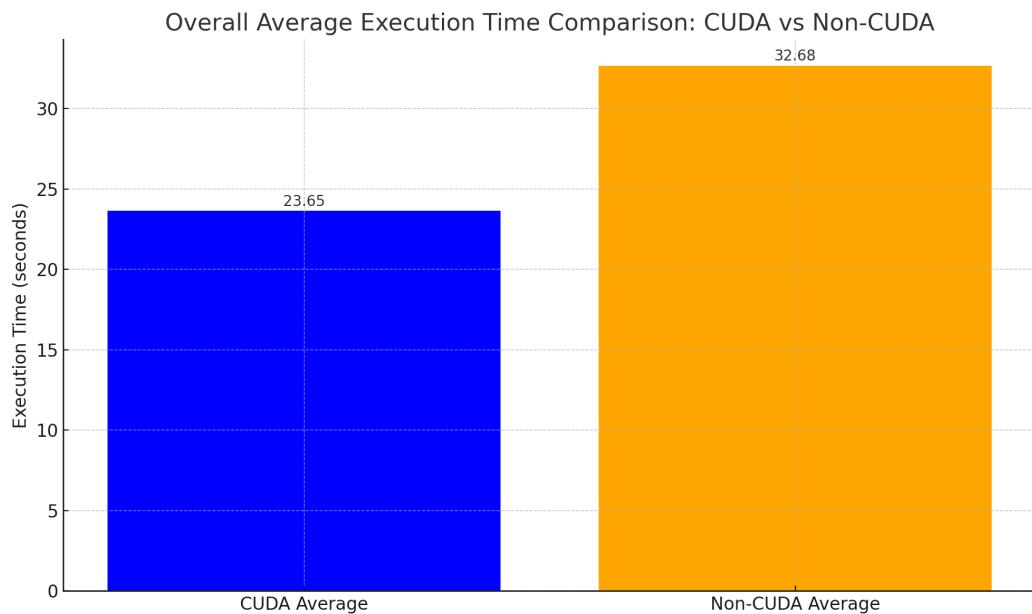
## =====Video2=====



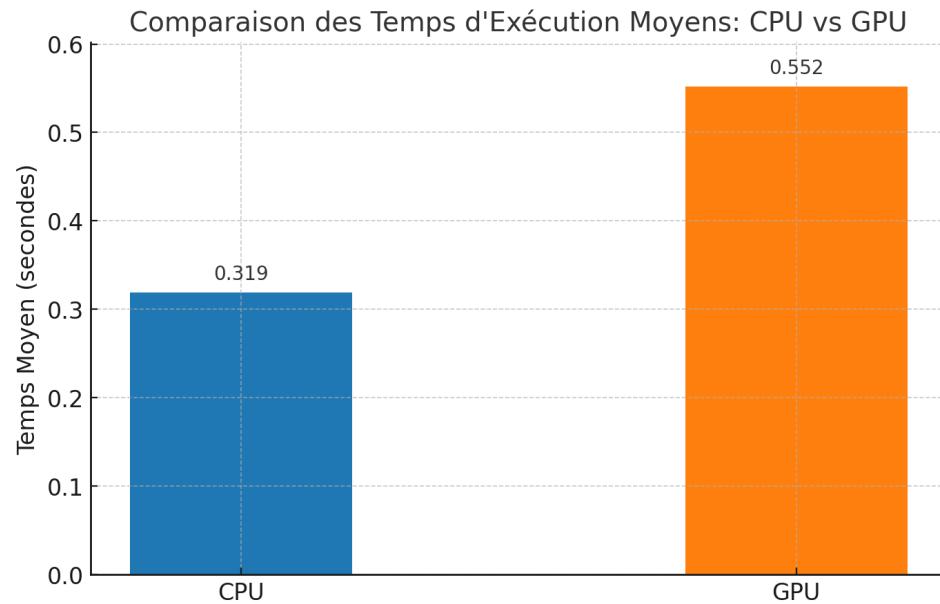


## ===== Video1 =====





## CuPy vs Numpy



---

# PANORAMA

---

## Image

Une image c'est grosse matrice de pixels.

```
[138 140 158]  
[136 138 156]  
[127 129 147]
```

Chaque ligne de la matrice représente un pixel donc numpy va etre important pour la suite

```
[Array([[[151, 127, ...ype=uint8), Array([[[151, 127, 89],  
> special variables  
> function variables  
> 00: Array([[[151, 127, 89],  
> 01: Array([[[151, 127, 89],  
> 02: Array([[[151, 127, 89],  
> 03: Array([[[152, 128, 90],  
> 04: Array([[[151, 127, 89],  
> 05: Array([[[152, 128, 90],  
> 06: Array([[[152, 128, 90],  
> 07: Array([[[152, 128, 88],  
> 08: Array([[[152, 128, 90],  
> 09: Array([[[152, 128, 90],  
> 10: Array([[[153, 129, 91],  
> 11: Array([[[153, 129, 91],  
> 12: Array([[[152, 128, 88],  
> 13: Array([[[154, 130, 90],  
> 14: Array([[[153, 129, 91],  
> 15: Array([[[153, 129, 91],  
> 16: Array([[[155, 129, 89],  
> 17: Array([[[154, 130, 90],  
> 18: Array([[[154, 130, 90],  
> 19: Array([[[154, 130, 90],  
> 20: Array([[[154, 130, 90],  
> 21: Array([[[154, 130, 90],  
> 22: Array([[[154, 130, 90],  
> 23: Array([[[154, 130, 90],  
Hold Alt key to switch to editor language hover
```

`cv2.imread()` Va lire l'image et la convertit en une matrice

## Algo de detection de point clé → ORB

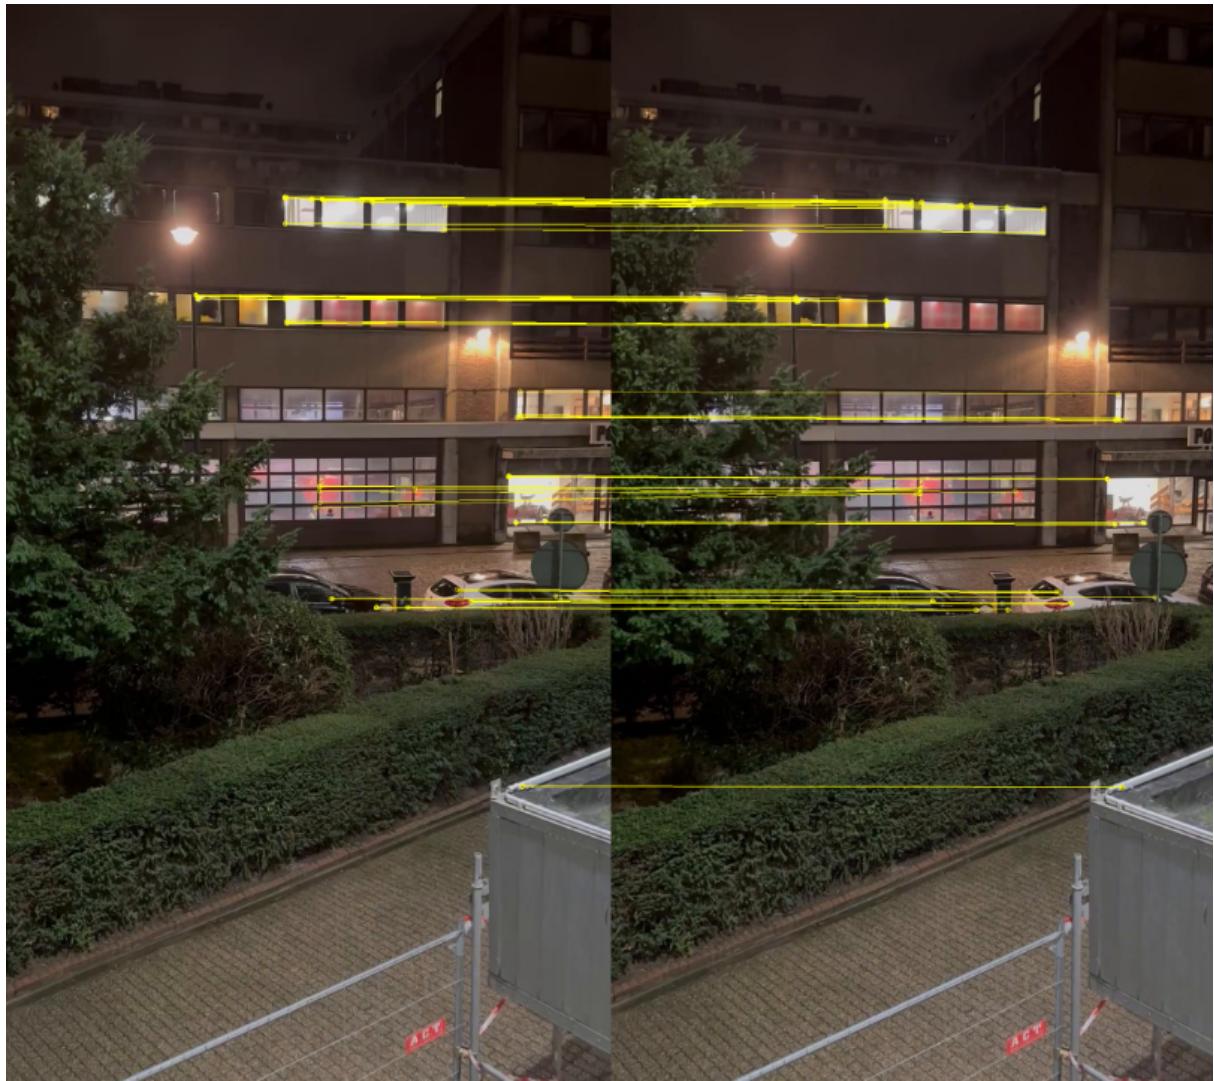
OpenCV: ORB (Oriented FAST and Rotated BRIEF)

In this chapter,

🔗 [https://docs.opencv.org/3.4/d1/d89/tutorial\\_py\\_orb.html](https://docs.opencv.org/3.4/d1/d89/tutorial_py_orb.html)

Trouver des régions ou les img match, et appliquer une transformation de perspective pourqu'elle soit positionner sur le même plan ou la première sera la cadre de ref et la 2em l'image déformé, de manière à aligné les caractéristique des 2 images

## Récupérer les données qui match



## Superposition

```
min_coords = list_of_points.min(axis=0).ravel() - 0.5 #j'aplati le tableau et je soustrait 0.5
```

